



**Turning job-search evidence into
better-informed decisions**

Ironhack AI Product Manager Final Project · Jose Manuel Lozano Barba



Meet Alex.

An active job seeker managing multiple applications at the same time.



Alex

Active job seeker

Multiple roles · Multiple decisions

CORE NEED

Know where to focus next.

Alex has the information.
But not always the clarity to act on it.



OPPORTUNITIES

Which roles are worth the effort?



APPLICATIONS

What should I tailor or change?




OUTCOMES

What might my results be telling me?



Research changed the product.

The evidence did not validate the original idea unchanged — it redirected it.



AI accelerated research synthesis + hypothesis organisation

Human decision changed the original concept instead of defending it



Competitors confirmed the need to differentiate.

The opportunity was not to copy existing workflow tools — it was to connect evidence with interpretation and next actions.


Tool	Strong in	What it proved	CareerPilot decision
 flowcv	CV and cover-letter creation	CV creation is already well served.	Do not compete on advanced visual CV building.
 huntr	Job capture, tracking, workflow and AI docs	Tracker + AI is already an established proposition.	Make tracking support interpretation, not define the product.
 teal	Integrated job-search workspace, resume tools and metrics	All-in-one workspace alone is not enough differentiation.	Focus on evidence-based decision support and next actions.

Strategic conclusion: CareerPilot should not be another CV builder, tracker or analytics dashboard.



Roadmap: focus first, expand after validation.

The roadmap protects the MVP from becoming a full job-search platform too early.



Build enough to validate → expand only when evidence supports it

 **AI accelerated** Rovo organised roadmap + backlog

 **Human decision** set scope, priorities and strategic boundaries



Human review is a product requirement.

One story turns the product principle into testable behaviour.

US-31

Review AI Recommendation

As an active job seeker, I want to review the context and reasoning behind an AI recommendation, so that I can decide whether it is relevant before taking any action.

[Strategy](#) → [backlog](#) → [implementation](#) → [demo](#)

ACCEPTANCE CRITERIA



Clear action



Evidence + context



Facts ≠ AI interpretation



Uncertainty visible



No automatic action

AI may suggest · the user must understand and decide



From evidence to a user-controlled recommendation

Watch the decision-support loop — not a feature tour.



CareerPilot suggests. The job seeker decides.



AI accelerated

Lovable accelerated implementation + iteration



Human decision

audited the build and kept the product exclusions



A working MVP is not a validated product.

The measurement framework asks whether users reach the experience — and perceive the recommendations as useful.

MVP NORTH STAR

Recommendation Usefulness Rate


Do users explicitly rate the recommendation as useful?

PROVISIONAL VALIDATION TARGET
≥70% usefulness · ≥50% feedback coverage

Capstone validation threshold · not an industry benchmark

MONITORING DASHBOARD

[Open dashboard ↗](#)




Current values are simulated demo data; the linked sheet is the live monitoring source.



Useful AI still needs boundaries.

Responsible AI is part of the product behaviour — not a policy slide added afterwards.




 **AI accelerated** drafting + consistency review of guardrails

 **Human decision** defined final boundaries and escalation rules



The next milestone is evidence, not more features.

The MVP is the beginning of validation — not the end of the product process.



Better evidence for decisions the job seeker still owns.

